

Goal

Deliver a production-ready relational database for PsychSync with: (1) a complete ERD, (2) schemas & migrations wired to FastAPI, (3) isolated dev/staging/prod environments, and (4) a repeatable workflow. Target DB: PostgreSQL 16.

0) Quick Decisions (stick to these unless there's a strong reason to change)

- **Engine:** PostgreSQL 16
 - **ORM:** SQLAlchemy 2.x (async) + Alembic
 - **Driver:** `asyncpg`
 - **IDs:** UUID v4 (database-generated default `gen_random_uuid()`)
 - **Timestamps:** `timestamptz` with `created_at` / `updated_at` triggers
 - **Soft delete:** `deleted_at` nullable
 - **Tenancy:** Org-scoped multi-tenant via foreign keys (no shared row data across orgs)
 - **Schema:** single `public` schema (keep it simple for v1)
 - **Naming:** snake_case for tables/columns; singular table names for entities with junction tables pluralized (e.g., `team_members`)
 - **Constraints:** NOT NULL + CHECKs + UNIQUE where relevant; `ON DELETE CASCADE` only on pure junctions
 - **Indexes:** compound indexes for hot queries noted below
-

1) Core Entities (MVP)

Identity & Access - `users` - platform accounts - `organizations` - customer accounts - `org_members` - user↔org membership (role at org level) - `teams` - teams belong to an organization - `team_members` - user↔team membership (role at team level) - `roles` (seeded: `owner`, `admin`, `member`, `viewer`) - `permissions` (coarse-grained, map to roles)

Assessment Framework - `frameworks` - e.g., MBTI, Big Five, DISC - `questions` - belongs to a framework - `question_options` - options for multiple choice - `assessments` - an assigned/taken assessment instance (user, framework, team, status) - `responses` - per question answer for an assessment - `scores` - normalized result per dimension (e.g., Big Five OCEAN)

Analytics & Optimization - `compatibility_rules` - rules used by the optimization engine - `team_metrics` - aggregated metrics per team & period

Operations - `invitations` - email invites to join org/team - `audit_logs` - security & data change logs - `notifications` - queued in-app/email events - `files` - uploaded artifacts (optional for v1) - `subscriptions` - (optional: Stripe integration in later phase)

Nice-to-have later: `segments`, `experiments`, `webhooks_outbox` (for reliable integrations), `idempotency_keys` for API writes.

2) ERD (Mermaid)

Copy/paste into docs/wiki to render. (Reflects MVP above.)

```
erDiagram
    organizations ||--o{ users : "owner (via org_members)"
    organizations ||--o{ org_members : has
    users ||--o{ org_members : has

    organizations ||--o{ teams : has
    teams ||--o{ team_members : has
    users ||--o{ team_members : has

    roles ||--o{ org_members : grants
    roles ||--o{ team_members : grants

    frameworks ||--o{ questions : has
    questions ||--o{ question_options : has

    frameworks ||--o{ assessments : assigned_as
    users ||--o{ assessments : takes
    teams ||--o{ assessments : context

    assessments ||--o{ responses : has
    assessments ||--o{ scores : produces

    organizations ||--o{ invitations : issues
    organizations ||--o{ audit_logs : records
    users ||--o{ audit_logs : actor
    teams ||--o{ team_metrics : aggregates
```

3) SQL DDL (PostgreSQL) – Seed Migrations

Create **initial migration** with these representative tables. You can expand with more columns as needed.

```
-- Enable extensions
CREATE EXTENSION IF NOT EXISTS pgcrypto;
CREATE EXTENSION IF NOT EXISTS citext;           -- for gen_random_uuid
                                                -- case-insensitive email/
                                                -- username
```

```

-- Timestamp trigger function
CREATE OR REPLACE FUNCTION set_updated_at()
RETURNS TRIGGER AS $$
BEGIN
    NEW.updated_at = NOW();
    RETURN NEW;
END;$$ LANGUAGE plpgsql;

-- organizations
CREATE TABLE organizations (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name TEXT NOT NULL,
    slug TEXT UNIQUE NOT NULL,
    created_at timestamp NOT NULL DEFAULT NOW(),
    updated_at timestamp NOT NULL DEFAULT NOW(),
    deleted_at timestamp
);
CREATE TRIGGER organizations_set_updated
BEFORE UPDATE ON organizations FOR EACH ROW EXECUTE PROCEDURE set_updated_at();

-- users
CREATE TABLE users (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    email CITEXT NOT NULL UNIQUE,
    password_hash TEXT NOT NULL,
    full_name TEXT,
    avatar_url TEXT,
    is_active BOOLEAN NOT NULL DEFAULT TRUE,
    created_at timestamp NOT NULL DEFAULT NOW(),
    updated_at timestamp NOT NULL DEFAULT NOW(),
    deleted_at timestamp
);
CREATE INDEX users_email_idx ON users (email);
CREATE TRIGGER users_set_updated
BEFORE UPDATE ON users FOR EACH ROW EXECUTE PROCEDURE set_updated_at();

-- roles (seed data in a later migration)
CREATE TABLE roles (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    scope TEXT NOT NULL CHECK (scope IN ('org','team')),
    code TEXT NOT NULL, -- e.g., 'owner','admin','member','viewer'
    UNIQUE(scope, code)
);

-- org_members
CREATE TABLE org_members (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```

org_id UUID NOT NULL REFERENCES organizations(id) ON DELETE CASCADE,
user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
role_id UUID NOT NULL REFERENCES roles(id),
created_at timestamp NOT NULL DEFAULT NOW(),
UNIQUE (org_id, user_id)
);
CREATE INDEX org_members_org_user_idx ON org_members (org_id, user_id);

-- teams
CREATE TABLE teams (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    org_id UUID NOT NULL REFERENCES organizations(id) ON DELETE CASCADE,
    name TEXT NOT NULL,
    description TEXT,
    created_at timestamp NOT NULL DEFAULT NOW(),
    updated_at timestamp NOT NULL DEFAULT NOW(),
    deleted_at timestamp,
    UNIQUE (org_id, name)
);
CREATE TRIGGER teams_set_updated
BEFORE UPDATE ON teams FOR EACH ROW EXECUTE PROCEDURE set_updated_at();
CREATE INDEX teams_org_idx ON teams (org_id);

-- team_members
CREATE TABLE team_members (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    team_id UUID NOT NULL REFERENCES teams(id) ON DELETE CASCADE,
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    role_id UUID NOT NULL REFERENCES roles(id),
    created_at timestamp NOT NULL DEFAULT NOW(),
    UNIQUE (team_id, user_id)
);
CREATE INDEX team_members_team_user_idx ON team_members (team_id, user_id);

-- frameworks
CREATE TABLE frameworks (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    code TEXT NOT NULL UNIQUE, -- 'MBTI','BIG5','DISC'
    name TEXT NOT NULL,
    created_at timestamp NOT NULL DEFAULT NOW(),
    updated_at timestamp NOT NULL DEFAULT NOW()
);
CREATE TRIGGER frameworks_set_updated
BEFORE UPDATE ON frameworks FOR EACH ROW EXECUTE PROCEDURE set_updated_at();

-- questions
CREATE TABLE questions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```

framework_id UUID NOT NULL REFERENCES frameworks(id) ON DELETE CASCADE,
code TEXT,           -- optional external code
body TEXT NOT NULL,
kind TEXT NOT NULL CHECK (kind IN ('single','multi','likert','text')),
position INT NOT NULL,
created_at timestamptz NOT NULL DEFAULT NOW()
);
CREATE INDEX questions_framework_pos_idx ON questions (framework_id, position);

-- question_options
CREATE TABLE question_options (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    question_id UUID NOT NULL REFERENCES questions(id) ON DELETE CASCADE,
    label TEXT NOT NULL,
    value TEXT NOT NULL, -- store raw value
    weight NUMERIC,      -- optional scoring weight
    position INT NOT NULL
);
CREATE INDEX question_options_q_pos_idx ON question_options (question_id,
position);

-- assessments (an instance of a user taking a framework)
CREATE TABLE assessments (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    org_id UUID NOT NULL REFERENCES organizations(id) ON DELETE CASCADE,
    team_id UUID REFERENCES teams(id) ON DELETE SET NULL,
    user_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    framework_id UUID NOT NULL REFERENCES frameworks(id) ON DELETE RESTRICT,
    status TEXT NOT NULL CHECK (status IN
('assigned','in_progress','completed','expired')),
    started_at timestamptz,
    completed_at timestamptz,
    created_at timestamptz NOT NULL DEFAULT NOW()
);
CREATE INDEX assessments_lookup_idx ON assessments (org_id, team_id, user_id,
framework_id);

-- responses
CREATE TABLE responses (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    assessment_id UUID NOT NULL REFERENCES assessments(id) ON DELETE CASCADE,
    question_id UUID NOT NULL REFERENCES questions(id) ON DELETE CASCADE,
    answer_json JSONB NOT NULL, -- supports multi/likert/text
    created_at timestamptz NOT NULL DEFAULT NOW(),
    UNIQUE (assessment_id, question_id)
);
CREATE INDEX responses_assessment_idx ON responses (assessment_id);

```

```

-- scores (normalized dimension scores)
CREATE TABLE scores (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    assessment_id UUID NOT NULL REFERENCES assessments(id) ON DELETE CASCADE,
    dimension TEXT NOT NULL,          -- e.g., 'O','C','E','A','N' or 'I/E'
    value NUMERIC NOT NULL,          -- 0..100 or normalized 0..1
    created_at timestamptz NOT NULL DEFAULT NOW(),
    UNIQUE (assessment_id, dimension)
);
CREATE INDEX scores_assessment_dim_idx ON scores (assessment_id, dimension);

-- invitations
CREATE TABLE invitations (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    org_id UUID NOT NULL REFERENCES organizations(id) ON DELETE CASCADE,
    email CITEXT NOT NULL,
    token TEXT NOT NULL UNIQUE,
    expires_at timestamptz NOT NULL,
    accepted_at timestamptz
);
CREATE INDEX invitations_org_email_idx ON invitations (org_id, email);

-- audit_logs
CREATE TABLE audit_logs (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    org_id UUID NOT NULL REFERENCES organizations(id) ON DELETE CASCADE,
    actor_user_id UUID REFERENCES users(id) ON DELETE SET NULL,
    action TEXT NOT NULL,           -- e.g., 'user.invited','assessment.completed'
    entity TEXT NOT NULL,          -- table name
    entity_id UUID,
    meta JSONB,
    created_at timestamptz NOT NULL DEFAULT NOW()
);
CREATE INDEX audit_logs_org_time_idx ON audit_logs (org_id, created_at DESC);

```

Seed roles (separate migration):

```

INSERT INTO roles (scope, code) VALUES
('org','owner'),('org','admin'),('org','member'),('org','viewer'),
('team','admin'),('team','member'),('team','viewer')
ON CONFLICT DO NOTHING;

```

4) Wire Up SQLAlchemy (Async) + Alembic (FastAPI)

Install deps

```
pip install "sqlalchemy[asyncio]" asyncpg alembic psycopg2-binary python-dotenv
```

Project structure (backend)

```
backend/
  app/
    db/
      __init__.py
      models.py
      session.py
    api/
    core/
    ...
  alembic/
  alembic.ini
```

app/db/session.py

```
from sqlalchemy.ext.asyncio import create_async_engine, async_sessionmaker
from sqlalchemy.orm import DeclarativeBase
import os

DATABASE_URL = os.getenv("DATABASE_URL") # e.g., postgresql+asyncpg://
                                         user:pass@localhost:5432/psychsync

class Base(DeclarativeBase):
    pass

engine = create_async_engine(DATABASE_URL, pool_pre_ping=True)
AsyncSessionLocal = async_sessionmaker(engine, expire_on_commit=False)
```

app/db/models.py (sample subset)

```
from sqlalchemy.orm import Mapped, mapped_column, relationship
from sqlalchemy import ForeignKey, text
from .session import Base
```

```

class Organization(Base):
    __tablename__ = "organizations"
    id: Mapped[str] = mapped_column(primary_key=True,
server_default=text("gen_random_uuid()"))
    name: Mapped[str]
    slug: Mapped[str]

class User(Base):
    __tablename__ = "users"
    id: Mapped[str] = mapped_column(primary_key=True,
server_default=text("gen_random_uuid()"))
    email: Mapped[str]
    password_hash: Mapped[str]

```

Initialize Alembic

`alembic init -t async alembic`

Edit `alembic.ini` - Set `script_location = alembic`

Edit `alembic/env.py` - load URL from env & use async engine

```

from logging.config import fileConfig
from sqlalchemy import pool
from sqlalchemy.engine import Connection
from alembic import context
import os

# Import models for autogenerate
from app.db.session import Base
from app.db import models # noqa

config = context.config
fileConfig(config.config_file_name)

target_metadata = Base.metadata

def get_url():
    return os.getenv("DATABASE_URL")

def run_migrations_offline():
    url = get_url()
    context.configure(url=url, target_metadata=target_metadata,
literal_binds=True, dialect_opts={"paramstyle": "named"})

```

```

    with context.begin_transaction():
        context.run_migrations()

def run_migrations_online():
    from sqlalchemy.ext.asyncio import async_engine_from_config
    connectable = async_engine_from_config(
        {"sqlalchemy.url": get_url()}, prefix="sqlalchemy.",
        poolclass=pool.NullPool
    )
    async def do_run_migrations():
        async with connectable.connect() as connection:
            await connection.run_sync(lambda conn:
context.configure(connection=conn, target_metadata=target_metadata))
                await connection.run_sync(lambda conn: context.begin_transaction())
                await connection.run_sync(lambda conn: context.run_migrations())
import asyncio
asyncio.run(do_run_migrations())

if context.is_offline_mode():
    run_migrations_offline()
else:
    run_migrations_online()

```

Create migrations

```

# 1) Revision for extensions & triggers (manual SQL)
alembic revision -m "extensions & triggers"
# paste SQL into the revision's upgrade() using op.execute(...)

# 2) Autogenerate tables from models
alembic revision --autogenerate -m "core tables"

# 3) Seed roles
alembic revision -m "seed roles"
# add op.execute(INSERT ...) in upgrade()

# Apply
alembic upgrade head

```

5) Environments

Env Vars (in each env)

```
DATABASE_URL=postgresql+asyncpg://psychsync:psychsync@localhost:5432/psychsync
DATABASE_URL_TEST=postgresql+asyncpg://psychsync:psychsync@localhost:5432/
psychsync_test
```

Docker Compose (dev)

```
version: '3.8'
services:
  db:
    image: postgres:16
    environment:
      POSTGRES_USER: psychsync
      POSTGRES_PASSWORD: psychsync
      POSTGRES_DB: psychsync
    ports: ["5432:5432"]
    volumes:
      - pgdata:/var/lib/postgresql/data
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U $$POSTGRES_USER"]
      interval: 5s
      timeout: 5s
      retries: 10
  volumes:
    pgdata: {}
```

Staging/Prod - Use managed Postgres (AWS RDS/Azure Flexible Server) - Private networking + TLS required
- Automated daily snapshots; PITR enabled - Separate DB per env (no shared data), distinct credentials

Migrations in CI/CD - Run `alembic upgrade head` on deploy (staging, then prod) - Block deploy if migrations fail - Keep a `down_revision` chain linear; squash old migrations periodically

6) Data & Access Patterns (index plan)

- `assessments_lookup_idx (org_id, team_id, user_id, framework_id)` - dashboard filters
 - `responses(assessment_id)` - loading answer sets
 - `scores(assessment_id, dimension)` - reading per-dimension quickly
 - `teams(org_id, name)` - unique per org; search by name within org
 - `audit_logs(org_id, created_at desc)` - recent activity feed
-

7) Security & Compliance Basics (v1)

- Store only necessary PII; use `citext` for emails
 - Hash passwords with **Argon2id** (app layer) & long random salts
 - Restrict by org: **every query** that reads/writes domain data must filter by `org_id`
 - Do not expose internal IDs in URLs without authorization checks
 - Enable **pg_stat_statements** in staging/prod for performance analysis
 - Periodic anonymized snapshots for staging (no raw PII)
-

8) Test Data & Seeding

Create a Python seed script to populate a dev org, users, teams, and a minimal Big Five framework with 5-10 questions to unblock UI work.

```
# scripts/seed.py (run with: uvicorn not needed; use Python)
import asyncio
from app.db.session import AsyncSessionLocal, engine
from app.db.models import Organization, User

async def main():
    async with AsyncSessionLocal() as s:
        acme = Organization(name="Acme Corp", slug="acme")
        s.add(acme)
        await s.flush()
        u = User(email="founder@acme.com", password_hash="dev_only",
                  full_name="Dev Founder")
        s.add(u)
        await s.commit()

asyncio.run(main())
```

9) Definition of Done (for this task)

- ERD checked into repo (`/docs/erd.mmd` + exported PNG)
 - Alembic migrations: extensions/triggers, core tables, seed roles
 - Local dev DB up via Docker; `alembic upgrade head` succeeds
 - FastAPI connects with async SQLAlchemy; healthcheck endpoint passes
 - Basic seed script creates org, user, team, framework, questions
 - Staging DB provisioned with credentials stored in secrets manager
 - CI job runs migrations automatically on staging
-

10) 16-Hour Work Plan (Backend Lead)

1. **Design** (3h): finalize entities & ERD, review with Tech Lead
 2. **Scaffold** (2h): SQLAlchemy base, session, Alembic init
 3. **Migrations** (4h): extensions/triggers + core tables + seeds
 4. **Env Setup** (2h): Docker Compose, .envs, Makefile targets
 5. **Wiring** (2h): FastAPI startup checks, connection lifecycle
 6. **Seeding & Smoke** (2h): seed script + CRUD sanity
 7. **Docs & Handover** (1h): README, ERD image, usage notes
-

11) Makefile (quality of life)

```
.PHONY: db-up db-down migrate upgrade seed

DB_URL=postgresql+asyncpg://psychsync:psychsync@localhost:5432/psychsync

db-up:
    docker compose up -d db
    sleep 3

migrate:
    alembic revision --autogenerate -m "$(m)"

upgrade:
    DATABASE_URL=$(DB_URL) alembic upgrade head

seed:
    python scripts/seed.py
```

12) Next Actions (execute in order)

- 1) docker compose up -d db
- 2) Set DATABASE_URL in .env and shell
- 3) alembic upgrade head
- 4) make seed (or python scripts/seed.py)
- 5) Run FastAPI; hit /health to verify DB connectivity
- 6) Export ERD PNG, add to repo, share in sprint review